

TRANSACTION PROCESSING

1. Transactions

A transaction is a unit of program execution that accesses and possibly updates various data items i.e., a sequence of operations that is regarded in a single logical operation is called a transaction. A transaction usually results from the execution of a user program written in a high-level data manipulation language or programming language (e.g. SQL, C etc) and is delimited by statements of the form '**begin transaction**' and '**end transaction**'. The transaction consists of all operations executed between the begin and end of the transaction.
A transaction can be modeled as a series of reads and writes.

Read (Q) – which transfers the data item Q from the database to a local buffer belonging to the transaction that executed the read operation.

Write (Q) – which transfers the data item Q from the local buffer of the transaction that executed the write back to database.

E.g.: - **Transaction Ti**- Transfers Rs.50 from account A to account **B**

Ti: read (A);
A=A-50;
Write (A);
Read (B);
B=B+50;
Write (B);

To ensure the integrity of the data, transactions should satisfy some properties called ACID Properties.

2. The ACID Properties

Atomicity all-or-nothing property

Either all operation of the transaction are reflected properly in the database or none are. Effect of an incomplete transaction should not be visible in DB

E.g.: If the system fails while executing **read (B)**, the balance of **A** will have been updated but that of **B** haven't. A Transaction is partially done, and the data is made inconsistent.

To ensure atomicity the db system keeps track (on disk) of the old values of any data on which a transaction perform a write and if the transaction does not complete its execution, the old values are restored to make it appear as though the transaction never executed. Ensuring atomicity is the responsibility of the db system specially, it is handled by a component called the **transaction manager component**.

Consistency

Database should be consistent before and after the execution of a transaction. Execution of a transaction in isolation (ie, with no other transaction executing concurrently) preserves the consistency of the db.

E.g. The consistency requirement in Ti is that the sum of A and B be unchanged by the execution of the transaction.

Ensuring consistency is the responsibility of the **application programmer** who codes the transaction. This task may be facilitated by automatic testing of **integrity constraints**.

Isolation

Each transaction is isolated from the effect of concurrent execution of other transactions. i.e., Transactions appear to be executed in serial order, even though there might be concurrent execution of other transactions.

E.g. consider the transaction Tj executing concurrently with Ti, If Tj reads A and B while executing Ti and computes A+B, it will observe an inconsistent value.

A solution to the problem of concurrently executing transactions is to execute transactions serially. i.e., one after the other. However, concurrent execution of transaction provides significant performance benefits. The isolation property ensures that the concurrent execution of transactions results in a system state that is equivalent to a state that could have been obtained had these transactions executed one at a time in some order. i.e., each transaction is unaware of other transactions executing concurrently in the system. Ensuring the isolation property is the responsibility of a component of the db system called in **concurrency controller**.

Durability

After a transaction completes successfully, the changes it has made to db persist, even if there are system failures.

Values of A and B may be in a memory buffer after the writes have been done before they are flushed to disk. A failure of the computers system may result in loss of data in main memory, but data written to disk are never lost.

We can guarantee durability by ensuring that either (1) updates carried out by the transaction have been written to disk before the transaction completes (2) Information about the updates carried out by the transaction and written to disk is sufficient to enable the db to reconstruct the updates when the db system is restarted after the failure. i.e., transactions are logged and are redone or undone after the system restarts. Ensuring durability is the responsibility of a component of the db system called **recovery management component**.

3. Transaction Operations

To support transaction processing, a DBMS usually provide three transaction operations: **Start, Commit, and Abort** for database programs. These operations are used in transaction implementations besides **Read, write and other arithmetic operations.**

Start: - Program tells the DBMS that is about to begin execution of a new transaction. The DBMS will have to make certain preparation for its execution.

Commit: - A Program tells the DBMS that the transaction has terminated normally and all of its effects should be made permanent. E.g., the disk should be updated to reflect changes made to the db. Execution of a transactions commit operation guarantees that abort of the transaction is not possible afterward.

Abort: - A program tells the DBMS that the transaction has terminated abnormally and all of its effects should be obliterated. Once the changes caused prior to transaction abortion is undone, we say that the transaction is rolled back to its initial states.

4. Schedule

A schedule is a list of operations (e.g. Read, Write, Start, Abort, Commit) from a set of transactions. The order of appearance of two operations of a particular transaction T in a schedule is the same as that in T.

E.g. Consider the transactions T_i and T_j

T_i -Transfers Rs.50 from a/c A to a/c B T_j -Transfer 10% of the balance from a/c A to account B

T_i : Read (A);
A=A-50;
Write (A);
Read (B);
B=B+50;
Write (B);
Write (B);

T_j : Read (A);
temp = A*0.1;
A=A-temp;
write (A)
read (B)
B=B+ temp;
write (B);

Serial Schedule

Each serial schedule consists of a sequence of operations from various transactions, where the operation belonging to one single transaction appear together in that schedule. For a set of n transactions ($T_1, T_2 \dots T_n$), there exist $n!$ different valid serial schedules.

T_n

Schedule 1 (T_i and T_j are executed one at a time in the order T_i followed by T_j)

T_i	T_j
Read (A)	
A=A-50	A
Write (A)	3000
Read (B)	B
B=B+50	3000
Write (B)	3000
	3000 + 500 = 3500
	3500
	3500 * 0.1 = 350
	3500 - 350 = 3150
	3150 + 500 = 3650
	3650
	3650 * 0.1 = 365
	3650 - 365 = 3285
	3285 + 500 = 3785
	3785
Read (A)	
Temp = A*0.1	A
Write (A)	3000
Read (B)	B
$\neg B=B+ temp$	3000 + 365 = 3365
Write (B)	3365

$A + B$ is consistent after the execution of the schedule

Schedule 2 Serial Schedule (T_j, T_i)

T_i	T_j
	Read (A)
	Temp = A*0.1
	Write (A)
	Read (B)
	$B=B + temp$
	Write (B)
Read (A)	
A=A-50	A
Write (A)	1000
Read (B)	B
B=B+50	1000 + 500 = 1500
Write	1500
Read (A)	
temp = A*0.1;	A
A=A-temp;	1000
write (A)	900
read (B)	B
B=B+ temp;	900 + 500 = 1400
write (B);	1400
Read (A)	
A=A-50	A
Write (A)	850
Read (B)	B
B=B+50	850 + 500 = 1350
Write	1350

$A + B$ is consistent after the execution of the schedule.

Concurrent schedule

If two transaction are running concurrently, the OS may execute one transaction for a little while then perform the context switch, execute the second transaction for some time, and then switch back to the first transaction for some time and so on. With multiple transactions, the CPU time is shared among all the transaction. In general it is not possible to predict exactly how many operations of a transaction will be executed before the CPU switches to another transaction. Thus the number of possible schedules for a set of n transaction is much larger than $n!$.

Schedule 3 (Concurrent schedule ---Ti and Tj are executed concurrently)

Ti	Tj
Read (A) A=A-50 Write (A)	
	Read (A) Temp = A*0.1 A=A-temp Write (A)
Read (B) B=B+50 Write (B)	Read (B) B=B +temp Write (B)
A + B is same before and after the execution of the schedule.	

Schedule 4 (Another concurrent schedule)

Ti	Tj
Read (A) A=A-50	
	Read (A) Temp = A*0.1 A=A-temp Write (A) Read (B)
Write (A) Read (B) B=B+50 Write (B)	B=B +temp Write (B)
Inconsistent state: A+B is not preserved by the execution of the schedule	

5. Concurrency controller

Concurrency controller controls the concurrent execution of transactions. If control of concurrent execution is left entirely to the OS, many possible schedules, including ones that leave the db in an inconsistent state such as schedule 4 are possible. It is the job of the db system to ensure that any schedule that gets executed will leave the db in a consistent state. The component of the db system that carries out this task is called concurrency controllers. We can ensure consistency of the db under

concurrent execution by making sure that any schedule that executes has the same effect as a schedule that could have occurred without any concurrent execution. i.e., the schedule should be equivalent to a serial schedule. Concurrency controller orders the db transaction operations so that the resulting schedule is equivalent to a serial schedule.

6. Serializability.

A schedule is said to be **serializable** if the result of executing that schedule is the same as the result of executing some serial schedule. So serializable schedule ensures consistency in concurrent execution of transactions. A **schedule S** having transactions Ti and Tj is serializable if S produces the same results as either $\langle T_i, T_j \rangle$ or $\langle T_j, T_i \rangle$.

Different conditions are satisfied for different kinds of schedule equivalence. They lead to the notions of **conflict serializability** and **view serializability**.

Conflict Serializability

Let T_i and T_j be two transactions with operations I_i and I_j respectively and Q be a data item in the database. If I_i and I_j refer to different data items the order of them is a schedule doesn't matter. We can swap I_i and I_j without affecting the results of any operation of the schedule. If I_i and I_j refer to the same data item Q then the order of them may matter.

There are 4 cases to consider:

$I_i = \text{read } (Q)$, $I_j = \text{read } (Q)$ the order of I_i and I_j doesn't matter.

$I_i = \text{read } (Q)$, $I_j = \text{write } (Q)$ the order of I_i and I_j matters

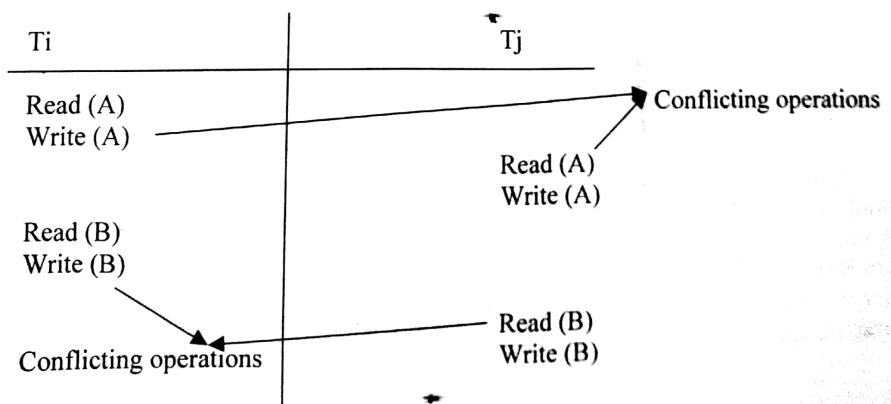
$I_i = \text{write } (Q)$, $I_j = \text{read } (Q)$ the order of I_i and I_j matters

$I_i = \text{write } (Q)$, $I_j = \text{write } (Q)$ the order of I_i and I_j matters

Two operations I_i and I_j are in conflict if they are operations by different transactions on the same item and at least one of them is a write. We can swap the order of adjacent operations in a schedule if they do not conflict with each other.

If a schedule S can be transformed to a schedule S' by a series of swaps of non-conflicting operation, S and S', are said to be conflict equivalent.

Schedule 3 (Considering only read and write operations)



Schedule 4 (After swapping non conflicting operation)

T _i	T _j
Read (A)	
Write (A)	
Read (B)	Read (A)
Write (B)	Write (A)
	Read (B) Write (B)

A schedule is Conflict Serializable if its conflict equivalent to a serial schedule.

E.g.: 1) After doing the following swapping of non conflicting operations in schedules3, the final result will be serial schedule 1

Swap Write (A) of T_j with Read (B) of T_i (We get schedule4)

Swap Read (B) of T_i with Read (A) of T_j.

Swap Write (B) of T_i with Write (A) of T_j

Swap Write (B) of T_i with Read (A) o T_j

The result is, schedule 1 (serial schedule)

T _i	T _j
Read (A)	
Write (A)	
Read(B)	Read(A)
Write (B)	Write (A)
	Read(B) Write (B)

Schedule 3 is conflict equivalent to schedule 1 which is a serial schedule; therefore schedule 3 is conflict serializable.

7

Testing serializability

Precedence Graphs

(V, E)

T_i - Vertices, V

T_i → T_j - Edges, E

There is edge T_i → T_j only if either of following:

1. T_i executes read(Q) before T_j executes write(Q)
2. T_i " " write(Q) " " T_j " " read(Q)
3. T_i " " write(Q) " " T_j " " write(Q)

No cycle in the Graph - schedule is conflict serializable

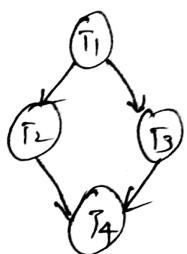
Cycle - not conflict serializable



E.g. 2) It is possible to have two schedules that produce the same outcome but that are not conflict equivalent.

Ti	Tj
Read (A) A=A-50 Write (A)	→ Read (B) B=B+10 Write (B)
Read (B) B=B+50 Write (B)	→ Read (A) A=A+10 Write (A)

Serializability Order can be obtained through topological sorting



$\langle T_1, T_2, T_3, T_4 \rangle$ or
 $\langle T_1, T_3, T_2, T_4 \rangle$.

The above schedule is not conflict equivalent to either of the serial schedules $\langle T_i, T_j \rangle$ or $\langle T_j, T_i \rangle$. But the final values of accounts A and B after the execution of either of them are the same. We can see from this example that there are less stringent definition of schedule equivalence than there are of conflict equivalence.

View Serializability

Consider two schedules S and S' where the same set of transactions participates in both schedules. S and S' are said to be view equivalent if the following three conditions are met.

1. For each data item Q, if T_i reads the initial value of Q in S then T_i must, in S', also, reads the initial value of Q.
2. For each data item Q, if T_i executes Read (Q) in S and that value was produced by T_j (if any) then T_i must, in S', also read the value of Q that was produced by T_j .
3. For each data item Q, the transaction (if any) that performs the final Write (Q) operation in S must perform the final write (Q) operation in S'.

E.g. Schedule 1 is not view equivalent to schedule 2
Schedule 1 is view equivalent to schedule 3

but that are

Schedule S is **View Serializable** if it is view equivalent to a serial schedule.
Eg. Schedule 3 is View Serializable since it is view equivalent to schedule which
is a serial schedule.

Every conflict Serializable schedule is View Serializable, but there are View Serializable schedules that are not Conflict Serializable.

Eg. 1) schedule 3 is view serializable as well as conflict serializable.

Ti	Tj	Tk
Read (Q)	Write (Q)	
Write (Q)		Write (Q)

This schedule is view serializable. It is view equivalent to serial schedule $\langle T_i, T_j, T_k \rangle$. But it is not conflict serializable.

Recoverability

Ti	Tj
read(A)	
write(A)	

read (A)
write(A)

If the schedule S consisting T_i and T_j is recoverable such that T_j is reading some value produced by T_i , the schedule is recoverable only if T_i commits before T_j .

Cascadeless Rollback

Ti	Tj	Tk
read(A) write(A) //commit	read(A) write(A) //commit	
		read (A)

A schedule S consisting of T_i, T_j and T_k . For ensuring cascadeless rollback, T_i should be committed before T_j and T_j should be committed before T_k .

Cascadeless Rollback schedule are Recoverable Schedule

CONCURRENCY CONTROL TECHNIQUES

There are two major types of concurrency control protocols:

- Lock-based protocols
- Timestamp-based protocols

1. Lock-based protocols

To ensure serializability and thus avoid violation of the isolation property, we access data items in a mutually exclusive manner. The idea is to allow transaction to access a data item only if it is holding a lock on it. Locks specify access mode of a data item. Two important kinds of locks are **shared and exclusive**.

Shared - If a transaction T_i has obtained a shared mode lock (denoted as S) on a data item Q, then T_i can read it but not write to it.

Exclusive - If a transaction T_i has obtained a exclusive-mode lock (denoted as X) on a data item Q, then T_i can both read it and write to it.

Lock	Read	Write
None	No	No
Shared	Yes	No
Exclusive	Yes	Yes

Many transactions can hold shared locks on a data item. If a transaction holds an exclusive lock on a data item no other transaction can hold any lock on it.

Lock Compatibility

Let A and B represent arbitrary lock modes. Suppose a transaction T_i requests a lock of mode A on item Q on which transaction T_j ($T_i \neq T_j$) currently holds a lock of mode B. If transaction T_i can be granted a lock on Q immediately, in spite of presence of the mode B lock, then we say mode A is compatible with mode B.

Lock compatibility matrix (between shared and exclusive locks) comp.

	S	X
S	True	False
X	False	false

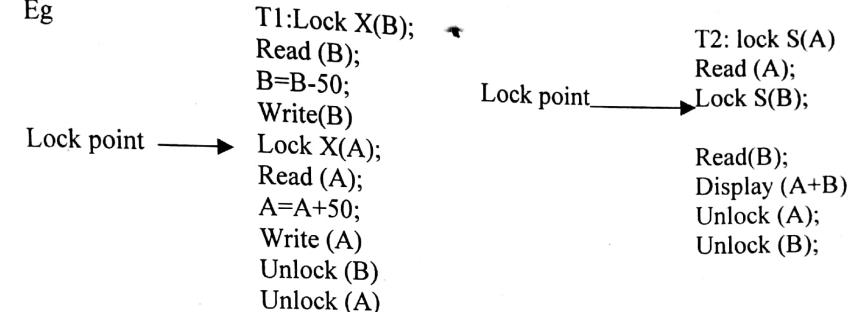
An element $comp(A, B)$ of the matrix has the value true if and only if mode A is compatible with mode B

Two Phase Locking Protocol

Each transaction issue lock and unlock requests in two phases.

1. **Growing phase** - a transaction may obtain locks, but may not release any lock.
2. **Shrinking phase** - A transaction may release locks, but may not obtain any new locks.

Eg



Initially, a transaction is in the growing phase. The transaction acquires locks as needed. Once the transaction releases a lock, it enters the shrinking phase, and it can issue no more lock requests. The point in the schedule where the transaction has obtained its final lock (the end of its growing phase) is called the lock point of the transaction.

We can show that the 2PL ensures conflict serializability. Transactions can be ordered according to their lock points – this ordering is in fact a **serialized** ordering for the transactions.

Note: (1) 2PL doesn't ensure freedom from deadlock.
(2) cascading roll back may occur under 2PL.

Variations of 2PL

Strict 2PL - The strict 2PL requires, that in addition to locking being 2φ all exclusive-mode locks taken by a transaction must be held until that transaction commits. This requirement ensures that any data written by an uncommitted transaction are locked in exclusive mode until the transaction commits, preventing any other transaction from reading the data. (Thus we can avoid cascading rollback)

Rigorous 2PL - is 2PL with the additional requirements that all locks be held until the transaction commits. It can be easily verified that with rigorous 2PL, transactions can be serialized in the order in which they commit. Most DBMS implement either strict or rigorous two-phase locking.

2PL

Example - Deadlock

LOCK-X(A)	
read(A)	
A = A - 50	
write(A)	
	LOCK-X(B)
	read(B)
	B = B + 50
	LOCK-X(A)
LOCK-X(B)	

Example - Cascading Roll back

T _i	T _j
LOCK-X(A)	
read(A)	
A = A - 50	
write(A)	
unlock(A)	
	LOCK-X(A)
	read(A)
	A = A + 50
	write(A)
	unlock(A)
---	---

Failure →

UP GRADATION

To avoid starvation -

e.g:-
 LOCK X(A1)
 read(A1)
 lock(X2)
 read(A2)
 lock S(A3)
 read(A3)
 write(A1)
 unlock(A1)

T_i

~~LOCK~~
~~read(A1)~~
~~read(A2)~~

unlock(A2)
 unlock(A3)

LOCK-S(A1) → starvation

read(A1)
 unlock(A1)

Protocol

1. If T_i issues read(Q), system issues lock-S(Q) followed by read(G)
2. When T_i issues write(Q), system issues upgrade(Q) if T_i in current holding a shared lock on Q. (This upgrade(Q) will wait if any other transaction T_j currently holding a shared lock on Q)
3. Else if T_i not currently holding a shared lock on Q, then issue lock-X(Q)
4. Release locks during commit operation where downgrading also possible

e.g:-

T _i	T _j
LOCK-S(A1)	LOCK-S(A1)
read(A2)	LOCK-S(A1)
---	read(A1)
read(A3)	LOCK-S(A1)
---	unlock(A1)

2. Timestamp-Based Protocols

In time stamp based protocols the transactions are ordered in advance using timestamps. With each transaction T_i in the system, we associate a unique fixed timestamp denoted by $t(T_i)$. The timestamp is assigned by the system before T_i starts execution. If a transaction T_i has been assigned timestamp $t(T_i)$, and a new transaction T_j enters the system, then $t(T_i) < t(T_j)$. Each data item we associate two time stamp values.

$t_w(Q)$ - denotes the largest timestamp of any transaction that executed $\text{Write}(Q)$ successfully.

$t_R(Q)$ - denotes the largest timestamp of any transaction that executed $\text{Read}(Q)$ successfully.

These two timestamps are updated whenever a new Read (Q) or Write (Q) instruction is executed.

The timestamps of the transactions determine the serializability order. Thus if $t(T_i) < t(T_j)$, then the system must ensure that the produced schedule is equivalent to a serial schedule in which T_i appears before T_j .

Timestamp-ordering Protocol

The timestamp ordering protocol ensures that any conflicting Read and Write operations are enclosed in timestamp order.

The protocol operates as follows:

1. Suppose T_i issues $\text{Read}(Q)$

- If $t(T_i) < t_w(Q)$, then T_i needs to read a value of Q that was already overwritten. Hence the read operation is rejected and T_i is rolled back.
- If $t(T_i) \geq t_w(Q)$ then the read operation is accepted and $t_R(Q)$ is set to the max of $t_R(Q)$ and $t(T_i)$

2. Suppose T_i issues $\text{Write}(Q)$

- $t(T_i) < t_R(Q)$ then the value of Q that T_i is producing was needed previously and the system assumed that the value would never be produced. Hence the write operation is rejected and T_i is rolled back.
- If $t(T_i) < t_w(Q)$ then T_i is attempting to write an ~~absolute~~ value of Q . Hence the write operation is rejected and T_i is rolled back.
- Otherwise, the write operation is executed and $t_w(Q)$ is set to $t(T_i)$

The timestamp ordering protocol ensures conflict serializability. This assertion follows from the fact that conflicting operations are processed in timestamp order. The protocol ensures freedom from deadlock, since no transaction ever waits.

e.g. T1: Read (A)
Write (A)
Read (B)
Write (B)

T2: Read (B)
Write (A)
Read (A)
Write (A)

Suppose $t(T_1) < t(T_2)$ and all t_0, t_1, t_2 and t_3 are smaller than $t(T_1)$

Schedule 1

T1	T2	$t_R(A)$	$t_w(A)$	$t_R(B)$	$t_w(B)$
Read (A)	-	t_0	t_1	t_2	t_3
Write (A)	-	$t(T_1)$	-	$t(T_1)$	-
Read (B)	-	-	-	$t(T_1)$	-
Write (B)	-	-	-	-	$t(T_1)$
Read (B)	-	-	-	$t(T_2)$	-
Write (B)	-	-	-	$t(T_2)$	-
Read (A)	$t(T_2)$	-	-	-	$t(T_2)$
Write (A)	-	-	$t(T_2)$	-	$t(T_2)$

Schedule 2

T1	T2	$t_R(A)$	$t_w(A)$	$t_R(B)$	$t_w(B)$
Read (A)	-	t_0	t_1	t_2	t_3
Write (A)	-	$t(T_1)$	$t(T_1)$	-	-
Read (B)	-	-	-	$t(T_2)$	-
Write(B)	-	-	-	-	$t(T_2)$
Read(A)	$t(T_2)$	-	-	-	-
Write(A)	-	-	$t(T_2)$	-	$t(T_2)$
Read(B)	-	-	-	$t(T_1)$	-
				$t(T_1) < t_w(B)$, T_1 is rolled back	

Thomas' Write Rule

This is a modified version of timestamp ordering protocol. The protocol rules for read operations remain unchanged. The protocol rules for write operations are slightly different from the timestamp ordering protocol.

Suppose T_i issues Write (Q)

1. If $t(T_i) < t_R(Q)$, then the value of Q that T_i is producing was previously needed, and it was assumed that the value would never be produced. Hence, the write operation is rejected and T_i is rolled back.
2. If $t(T_i) < t_w(Q)$, then T_i is attempting to write an obsolete value of Q. Hence this write operation can be ignored.
3. Otherwise, the write operation is executed and $t_w(Q)$ is set to $t(T_i)$

eg:

		Timestamp ordering protocol		Thomas write Rule	
T1	T2	$t_R(Q)$	$t_w(Q)$	$t_R(Q)$	$t_w(Q)$
Read (Q)		$t(T_1)$	-	$t(T_1)$	-
	Write (Q)	$t(T_1)$	$t(T_2)$	$t(T_1)$	$t(T_2)$
Write (Q)	→ Abort(), T_1 is rolled back			$t(T_1)$	$t(T_2)$

The roll back of T_i is required by time stamp ordering protocol. This is unnecessary because T_2 has already written Q, the value that T_1 is attempting to write is one that will never need to be read.

Any transaction T_i with $t(T_i) \leq t(T_2)$ that attempts a Read (Q) will be rolled back since $t(T_i) < t_w(Q)$.

Any transaction T_j with $t(T_j) > t(T_2)$ must read the value of Q written by T_2 , rather than by T_1 .

Advantage. Ensures view serializability. The schedule in the above example is not conflict serializable and this is not possible under any of 2φ locking, tree protocol or the timestamp ordering protocol. Under Thomas write rule the write (Q) operation of T_1 would be ignored. The result is a schedule that is view equivalent to the serial schedule $\langle T_1, T_2 \rangle$.